



How To Find Arbitrage Opportunities In Python

Hi there!

Let's solve an interesting programming interview problem: how to find an arbitrage. Here's the question.

Suppose you are given a table of currency exchange rates, represented as a 2D array. Determine whether there is a possible arbitrage: that is, whether there is some sequence of trades you can make, starting with some amount A of any currency, so that you can end up with some amount greater than A of that currency.

There are no transaction costs and you can trade fractional quantities.

How do we solve this?

Before continuing, you should take some time to try to solve it on your own!

We can model the currencies and the exchange rates as a graph, where the nodes are the currencies and the edges are the exchange rates between each currency. Since our table is complete, the graph is also complete. Then, to solve this problem, we need to find a cycle whose edge weights product is greater than 1.

This seems hard to do faster than brute force, so let's try to reduce it down to a problem we already know we can solve faster than brute force. (Scroll down if you want to see exactly how slow brute force is!)

Here's a hint: $\log(a * b) = \log(a) + \log(b)$.

If we take the negative log of the edge weights, the problem of finding a cumulative product that's greater than 1 turns into the problem of finding a negative sum cycle!

For example, say we have a weighted edge path $a \rightarrow b \rightarrow c \rightarrow d$. Since we want to see if $a * b * c * d > 1$, we'll take the negative log of each edge:

$-\log(a) \rightarrow -\log(b) \rightarrow -\log(c) \rightarrow -\log(d)$.

The total cost of this path will then be

$-(\log(a) + \log(b) + \log(c) + \log(d)) = -\log(a * b * c * d)$.

$-\log(x) < 0$ if x is greater than 1, so that's why if we have a negative cost cycle, it means that the product of the weighted edges is bigger than 1.

The Bellman-Ford algorithm can detect negative cycles. So if we run Bellman-Ford on our graph and discover one, then that means its corresponding edge weights multiply out to more than 1, and thus we can perform an arbitrage.

As a refresher, the Bellman-Ford algorithm is commonly used to find the shortest path between a source vertex and each of the other vertices. If the graph contains a negative cycle, however, it can detect it and throw an exception (or, in our case, return true). The main idea of Bellman-Ford is this:

Since the longest path in any graph has at most $|V| - 1$ edges, if we take all the direct edges from our source node, then we have all the one-edged shortest paths; once we take edges from there, we have all the two-edged shortest paths; all the way until $|V| - 1$ sized paths.

If, after $|V| - 1$ iterations of this, we can still find a smaller path, then there must be a negative cycle in the graph. We can start our algorithm on any vertex in our graph – since our graph is connected (by virtue of it being complete), then if there's a negative cycle in our graph, we'll find it.

```
from math import log
```

PYTHON

```
def arbitrage(table):
    transformed_graph = [[-log(edge) for edge in row] for row in graph]

    # Pick any source vertex -- we can run Bellman-Ford from any vertex and
    # get the right result
    source = 0
    n = len(transformed_graph)
    min_dist = [float('inf')] * n

    min_dist[source] = 0

    # Relax edges |V| - 1 times
    for i in range(n - 1):
        for v in range(n):
            for w in range(n):
                if min_dist[w] > min_dist[v] + transformed_graph[v][w]:
                    min_dist[w] = min_dist[v] + transformed_graph[v][w]

    # If we can still relax edges, then we have a negative cycle
    for v in range(n):
        for w in range(n):
            if min_dist[w] > min_dist[v] + transformed_graph[v][w]:
                return True

    return False
```

Because of the triply-nested for loop, this runs in $O(N^3)$ time, and because we're creating a transformed copy of the old graph we're using $O(N^2)$ space.

P.S: How slow is brute force? Well, we would have to check every single path in the graph that starts and ends at the same node, and our graph is complete. For each node, since our graph is connected, we can have every single permutation of any number of vertices up to $N - 1$, so brute force would be on the order of $O(N!)$ time!

I hope you enjoyed this interview question! Are you interviewing for programming jobs, or do you just enjoy fun programming questions? Check out our newsletter [Daily Coding Problem](#) to get a question in your inbox every day.

Ace your programming interview

Get a coding problem every day in your inbox!

Subscribe

PREV

NEXT